

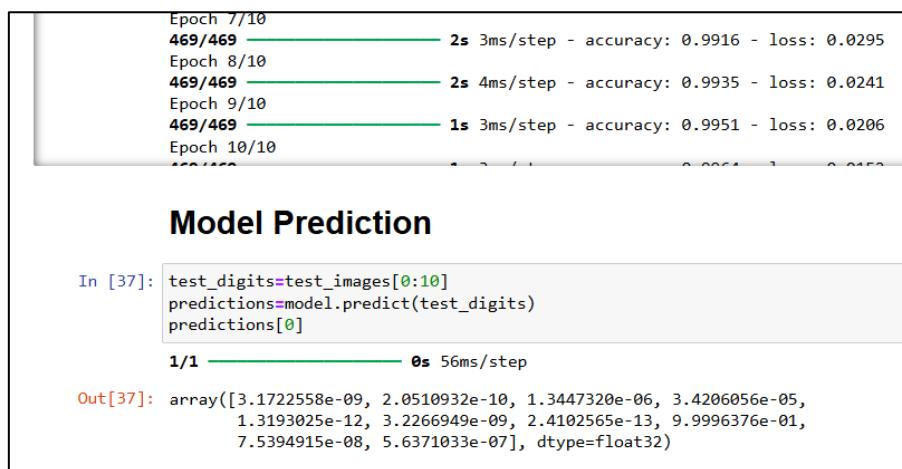
## Assignment-1

1. Create a neural network with one dense layer for MNIST dataset. Modify the network to achieve better accuracy by increasing the number of hidden layers and number of neurons.

**With one dense layer:**

```
from tensorflow import keras
from tensorflow.keras import layers
model=keras.Sequential([
    layers.Dense(256,activation='relu'),
    layers.Dense(10,activation='softmax')
])
```

**Output:**



The screenshot shows a Jupyter Notebook interface. At the top, there is a terminal-like output window displaying training logs for 10 epochs (7/10 to 10/10) on a dataset of 469 samples. Below this is a section titled "Model Prediction". In the "In [37]" cell, the code `test\_digits = test\_images[0:10]` and `predictions = model.predict(test\_digits)` is shown, along with the result `predictions[0]`. The "Out[37]" cell displays the predicted array: `array([3.1722558e-09, 2.0510932e-10, 1.3447320e-06, 3.4206056e-05, 1.3193025e-12, 3.2266949e-09, 2.4102565e-13, 9.9996376e-01, 7.5394915e-08, 5.6371033e-07], dtype=float32)`.

**After modifying:**

```
from tensorflow import keras
from tensorflow.keras import layers
model=keras.Sequential([
    layers.Dense(256,activation='relu'),
    layers.Dense(512,activation='relu'),
    layers.Dense(512,activation='relu'),
    layers.Dense(10,activation='softmax')
])
```

## Output:

```

469/469 ━━━━━━━━ 3s 7ms/step - accuracy: 0.9355 - loss: 0.2137
Epoch 6/10
469/469 ━━━━━━ 4s 9ms/step - accuracy: 0.9463 - loss: 0.1746
Epoch 7/10
469/469 ━━━━━━ 4s 9ms/step - accuracy: 0.9542 - loss: 0.1493
Epoch 8/10
469/469 ━━━━━━ 4s 9ms/step - accuracy: 0.9615 - loss: 0.1253
Epoch 9/10
469/469 ━━━━━━ 4s 9ms/step - accuracy: 0.9654 - loss: 0.1124
Epoch 10/10
469/469 ━━━━━━ 4s 9ms/step - accuracy: 0.9690 - loss: 0.1001

Out[44]: <keras.src.callbacks.history.History at 0x19d9f2dfbb0>

In [45]: test_digit=test_images[0:10]
predictions=model.predict(test_digit)
predictions[0]

1/1 ━━━━ 0s 69ms/step

Out[45]: array([8.1122249e-07, 1.1931706e-06, 2.7570373e-04, 3.6685797e-05,
   9.2656933e-08, 1.2184032e-07, 1.0611640e-12, 9.9966109e-01,
   2.3558825e-07, 2.4059318e-05], dtype=float32)

In [46]: predictions[0][7]

Out[46]: 0.9996611

```

## 2. Explain the maths of a MLP with loss function as Cross entropy (negative log likelihood) for softmax classifier. {Take hint from class notes for regression MLP}

A **Multi-layer Perceptron** is a type of neural network designed for supervised learning tasks such as regression and classification. It consists of an input layer, one or more hidden layers for feature extraction and an output layer that provides predictions and classified output. It also contains parameters weights and biases that are learned during training. Mathematics for MLP includes Forward Propagation, Activation Functions, Back Propagation and Loss Functions.

### 1. Forward Pass in an MLP

A Multilayer Perceptron (MLP) consists of multiple layers of neurons. Each layer transforms its input using weights, biases, and activation functions. Let's assume a network with:

- Input:  $x \in \mathbb{R}^d$  (where  $d$  is the input dimensionality).
- Hidden Layers: Each hidden layer applies a linear transformation followed by a non-linear activation function (e.g., ReLU).
- Output Layer: A final linear transformation followed by the softmax function.

### Hidden Layer Transformation

For a hidden layer  $h$ , the output is:

$$h = f(Wh + bh).$$

### Output Layer Transformation

The output layer computes logits  $z$ :

$$z = W_o h + b_o$$

### Softmax Function

The softmax function converts logits into probabilities:

$$\hat{y}_i = \exp(z_i) / \sum_j \exp(z_j)$$

Where  $\hat{y}_i$  is the predicted probability for class i.

## 2. Cross-Entropy Loss Function

The cross-entropy loss quantifies the difference between the predicted probability distribution and the true distribution. For a single example, the loss is:

$$L = -\sum_i y_i \log(\hat{y}_i)$$

Where:

- $y_i$  is the true label in one-hot encoding (1 for the correct class, 0 otherwise).
- $\hat{y}_i$  is the predicted probability for class i.

For a dataset with n samples:

$$L_{total} = (1/n) \sum_k L^{(k)}$$

Where  $L^{(k)}$  is the loss for the k<sup>th</sup> sample.

## 3. Gradient Computation (Backpropagation)

### Gradient of Softmax Function

The derivative of the softmax function with respect to the logits  $z_i$  is:

$$\partial \hat{y}_i / \partial z_j = \hat{y}_i (\delta_{ij} - \hat{y}_j)$$

Where  $\delta_{ij}$  is the Kronecker delta ( $\delta_{ij} = 1$  if  $i = j$ , else 0).

### Gradient of Cross-Entropy Loss

For the loss function:

$$\partial L / \partial z_i = \hat{y}_i - y_i$$

### Backpropagation Through Layers

1. Output Layer:

$$\partial L / \partial W_o = (\partial L / \partial z) h^T, \partial L / \partial b_o = \partial L / \partial z$$

2. Hidden Layers:

$$\partial L / \partial Wh = (\partial L / \partial h) x^T, \partial L / \partial bh = \partial L / \partial h$$

Where  $\partial L / \partial h$  depends on the activation function.

### 3. What is Vanishing gradient and Exploding gradient problem in neural networks.

#### **Vanishing Gradient Problem**

The vanishing gradient problem occurs during the training of deep neural networks when gradients of the loss function with respect to earlier layers become extremely small. This makes it difficult for the weights of these layers to update effectively, leading to slow or stalled learning.

Mathematically:

- During backpropagation, gradients are calculated using the chain rule:  

$$\frac{\partial L}{\partial W_1} = (\frac{\partial L}{\partial a_n}) (\frac{\partial a_n}{\partial a_{n-1}}) \dots (\frac{\partial a_2}{\partial a_1}) (\frac{\partial a_1}{\partial W_1}).$$
- If the activation functions have derivatives less than 1 (e.g., sigmoid or tanh), repeated multiplications can shrink the gradients exponentially as they propagate backward through layers
- As a result, the weights in earlier layers receive minimal updates.

Consequences:

- Slow convergence or no convergence during training.
- Poor performance of the model due to insufficient learning in early layers.

#### **Exploding Gradient Problem**

The exploding gradient problem occurs when gradients during backpropagation grow exponentially as they propagate backward through the layers. This results in extremely large updates to weights, which can destabilize the model.

Mathematically:

- During backpropagation, if activation functions or weights have large values, repeated multiplications can cause gradients to increase exponentially as they propagate backward through layers.
- This can lead to very large weight updates and numerical instability.

# Import libraries, load and transform data

In [1]:

```
| pip install -U -q evaluate transformers datasets >=2.14.5 accelerate >=0.27 mlflow
```

In [2]:

In [3]:

```
from PIL import ImageFile  
  
ImageFile.LOAD_TRUNCATED_IMAGES = True
```

In [4]:

```
image_dict = {}

from pathlib import Path
from tqdm import tqdm
import os

file_names = []
labels = []

for file in sorted((Path('/kaggle/input/cifake-real-and-ai-generated-synthetic-im')):
    label = str(file).split('/')[-2]
    labels.append(label)
    file_names.append(str(file))

print(len(file_names), len(labels))

df = pd.DataFrame.from_dict({"image": file_names, "label": labels})
print(df.shape)
```

120000 120000  
(120000, 2)

In [5]: df.head()

Out[5]:

	image	label
0	/kaggle/input/cifake-real-and-ai-generated-syn...	FAKE
1	/kaggle/input/cifake-real-and-ai-generated-syn...	FAKE
2	/kaggle/input/cifake-real-and-ai-generated-syn...	FAKE
3	/kaggle/input/cifake-real-and-ai-generated-syn...	FAKE
4	/kaggle/input/cifake-real-and-ai-generated-syn...	FAKE

In [6]: df['label'].unique()

Out[6]: array(['FAKE', 'REAL'], dtype=object)

```
In [7]: y = df[['label']]
df = df.drop(['label'], axis=1)
ros = RandomOverSampler(random_state=83)
df, y_resampled = ros.fit_resample(df, y)
del y
df['label'] = y_resampled
del y_resampled
gc.collect()
print(df.shape)
```

(120000, 2)

```
In [8]: dataset = Dataset.from_pandas(df).cast_column("image", Image())
```

```
In [9]: dataset[0]["image"]
```

Out[9]:



```
In [10]: labels_subset = labels[:5]
print(labels_subset)
```

['FAKE', 'FAKE', 'FAKE', 'FAKE', 'FAKE']

```
In [11]: labels_list = ['REAL', 'FAKE']
label2id, id2label = dict(), dict()
for i, label in enumerate(labels_list):
    label2id[label] = i
    id2label[i] = label

print("Mapping of IDs to Labels:", id2label, '\n')
print("Mapping of Labels to IDs:", label2id)
```

Mapping of IDs to Labels: {0: 'REAL', 1: 'FAKE'}

Mapping of Labels to IDs: {'REAL': 0, 'FAKE': 1}

```
In [18]: !pip uninstall -y ipywidgets
```

```
Found existing installation: ipywidgets 8.1.5
Uninstalling ipywidgets-8.1.5:
  Successfully uninstalled ipywidgets-8.1.5
```

In [19]: `!pip install ipywidgets`

```
Collecting ipywidgets
  Downloading ipywidgets-8.1.5-py3-none-any.whl (139 kB)
    139.8/139.8 kB 6.0 MB/s eta 0:00:0
0
Requirement already satisfied: comm>=0.1.3 in /opt/conda/lib/python3.10/site-packages (from ipywidgets) (0.1.3)
Requirement already satisfied: ipython>=6.1.0 in /opt/conda/lib/python3.10/site-packages (from ipywidgets) (8.14.0)
Requirement already satisfied: traitlets>=4.3.1 in /opt/conda/lib/python3.10/site-packages (from ipywidgets) (5.9.0)
Requirement already satisfied: widgetsnbextension~=4.0.12 in /opt/conda/lib/python3.10/site-packages (from ipywidgets) (4.0.13)
Requirement already satisfied: jupyterlab-widgets~=3.0.12 in /opt/conda/lib/python3.10/site-packages (from ipywidgets) (3.0.13)
Requirement already satisfied: backcall in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (0.2.0)
Requirement already satisfied: decorator in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (5.1.1)
Requirement already satisfied: jedi>=0.16 in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (0.18.2)
Requirement already satisfied: matplotlib-inline in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (0.1.6)
Requirement already satisfied: pickleshare in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (0.7.5)
Requirement already satisfied: prompt-toolkit!=3.0.37,<3.1.0,>=3.0.30 in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (3.0.38)
Requirement already satisfied: pygments>=2.4.0 in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (2.15.1)
Requirement already satisfied: stack-data in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (0.6.2)
Requirement already satisfied: pexpect>4.3 in /opt/conda/lib/python3.10/site-packages (from ipython>=6.1.0->ipywidgets) (4.8.0)
Requirement already satisfied: parso<0.9.0,>=0.8.0 in /opt/conda/lib/python3.10/site-packages (from jedi>=0.16->ipython>=6.1.0->ipywidgets) (0.8.3)
Requirement already satisfied: ptyprocess>=0.5 in /opt/conda/lib/python3.10/site-packages (from pexpect>4.3->ipython>=6.1.0->ipywidgets) (0.7.0)
Requirement already satisfied: wcwidth in /opt/conda/lib/python3.10/site-packages (from prompt-toolkit!=3.0.37,<3.1.0,>=3.0.30->ipython>=6.1.0->ipywidgets) (0.2.6)
Requirement already satisfied: executing>=1.2.0 in /opt/conda/lib/python3.10/site-packages (from stack-data->ipython>=6.1.0->ipywidgets) (1.2.0)
Requirement already satisfied: asttokens>=2.1.0 in /opt/conda/lib/python3.10/site-packages (from stack-data->ipython>=6.1.0->ipywidgets) (2.2.1)
Requirement already satisfied: pure-eval in /opt/conda/lib/python3.10/site-packages (from stack-data->ipython>=6.1.0->ipywidgets) (0.2.2)
Requirement already satisfied: six in /opt/conda/lib/python3.10/site-packages (from asttokens>=2.1.0->stack-data->ipython>=6.1.0->ipywidgets) (1.16.0)
Installing collected packages: ipywidgets
Successfully installed ipywidgets-8.1.5
```

```
In [21]: !pip list | grep ipywidgets
```

```
ipywidgets           8.1.5
```

```
In [22]: !jupyter nbextension enable --py widgetsnbextension --sys-prefix
```

```
Config option `kernel_spec_manager_class` not recognized by `EnableNBExtensionApp`.
```

```
Enabling notebook extension jupyter-js-widgets/extension...
```

```
- Validating: OK
```

```
In [23]: from datasets import ClassLabel

# Define ClassLabels
ClassLabels = ClassLabel(num_classes=len(labels_list), names=labels_list)

# Function to map Label names to IDs
def map_label2id(example):
    example['label'] = ClassLabels.str2int(example['label'])
    return example

# Apply mapping
dataset = dataset.map(map_label2id, batched=True)

# Cast the label column
dataset = dataset.cast_column('label', ClassLabels)

# Apply train-test split only to the 'train' split if it exists
if 'train' in dataset:
    dataset_split = dataset['train'].train_test_split(test_size=0.4, shuffle=True)
    dataset['train'] = dataset_split['train']
    dataset['test'] = dataset_split['test']
else:
    dataset = dataset.train_test_split(test_size=0.4, shuffle=True, stratify_by_col)

# Assign train and test datasets
train_data = dataset['train']
test_data = dataset['test']
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
In [24]: model_str = "dima806/ai_vs_real_image_detection"
processor = ViTImageProcessor.from_pretrained(model_str)
image_mean, image_std = processor.image_mean, processor.image_std
size = processor.size["height"]
print("Size: ", size)
normalize = Normalize(mean=image_mean, std=image_std)
_train_transforms = Compose(
    [
        Resize((size, size)),
        RandomRotation(90),
        RandomAdjustSharpness(2),
        ToTensor(),
        normalize
    ]
)

_val_transforms = Compose(
    [
        Resize((size, size)),
        ToTensor(),
        normalize
    ]
)

def train_transforms(examples):
    examples['pixel_values'] = [_train_transforms(image.convert("RGB")) for image in examples]
    return examples

def val_transforms(examples):
    examples['pixel_values'] = [_val_transforms(image.convert("RGB")) for image in examples]
    return examples
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

Size: 224

```
In [25]: train_data.set_transform(train_transforms)
test_data.set_transform(val_transforms)
```

```
In [26]: def collate_fn(examples):
    pixel_values = torch.stack([example["pixel_values"] for example in examples])
    labels = torch.tensor([example['label'] for example in examples])
    return {"pixel_values": pixel_values, "labels": labels}
```

## Load, train, and evaluate model

```
In [27]: model = ViTForImageClassification.from_pretrained(model_str, num_labels=len(label2id))
model.config.id2label = id2label
model.config.label2id = label2id
print(model.num_parameters(only_trainable=True) / 1e6)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

85.800194

```
In [29]: accuracy = evaluate.load("accuracy")
def compute_metrics(eval_pred):

    predictions = eval_pred.predictions
    label_ids = eval_pred.label_ids
    predicted_labels = predictions.argmax(axis=1)
    acc_score = accuracy.compute(predictions=predicted_labels, references=label_ids)
    return {
        "accuracy": acc_score
    }
```

```
In [30]: metric_name = "accuracy"
model_name = "ai_vs_real_image_detection"
num_train_epochs = 2
args = TrainingArguments(
    output_dir=model_name,
    logging_dir='./logs',
    evaluation_strategy="epoch",
    learning_rate=1e-6,
    per_device_train_batch_size=64,
    per_device_eval_batch_size=32,
    num_train_epochs=num_train_epochs,
    weight_decay=0.02,
    warmup_steps=50,
    remove_unused_columns=False,
    save_strategy='epoch',
    load_best_model_at_end=True,
    save_total_limit=1,
    report_to="none"
)
```

```
In [31]: trainer = Trainer(
    model,
    args,
    train_dataset=train_data,
    eval_dataset=test_data,
    data_collator=collate_fn,
    compute_metrics=compute_metrics,
    tokenizer=processor,
)
```

```
In [32]: trainer.evaluate()
```

[324/324 15:11]

```
Out[32]: {'eval_loss': 0.043403297662734985,
          'eval_model_preparation_time': 0.0035,
          'eval_accuracy': 0.9858217592592593,
          'eval_runtime': 141.4132,
          'eval_samples_per_second': 73.317,
          'eval_steps_per_second': 2.291}
```

```
In [33]: trainer.train()
```

[486/486 15:22, Epoch 2/2]

Epoch	Training Loss	Validation Loss	Model Preparation Time	Accuracy
1	No log	0.043228	0.003500	0.986400
2	No log	0.042479	0.003500	0.985436

```
Out[33]: TrainOutput(global_step=486, training_loss=0.034261570055298354, metrics={'train_runtime': 925.3218, 'train_samples_per_second': 33.614, 'train_steps_per_second': 0.525, 'total_flos': 2.4103108449722696e+18, 'train_loss': 0.034261570055298354, 'epoch': 2.0})
```

```
In [34]: trainer.evaluate()
```

```
Out[34]: {'eval_loss': 0.04247914254665375,
          'eval_model_preparation_time': 0.0035,
          'eval_accuracy': 0.9854359567901234,
          'eval_runtime': 91.5588,
          'eval_samples_per_second': 113.239,
          'eval_steps_per_second': 3.539,
          'epoch': 2.0}
```

```
In [35]: outputs = trainer.predict(test_data)
print(outputs.metrics)
```

```
{'test_loss': 0.04247914254665375, 'test_model_preparation_time': 0.0035, 'test_accuracy': 0.9854359567901234, 'test_runtime': 84.3297, 'test_samples_per_second': 122.946, 'test_steps_per_second': 3.842}
```

```
In [36]: y_true = outputs.label_ids
y_pred = outputs.predictions.argmax(1)

def plot_confusion_matrix(cm, classes, title='Confusion Matrix', cmap=plt.cm.Blue
    plt.figure(figsize=figsize)
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    fmt = '.0f'
    thresh = cm.max() / 2.0
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt), horizontalalignment="center", color="white" if cm[i, j] > thresh else "black")

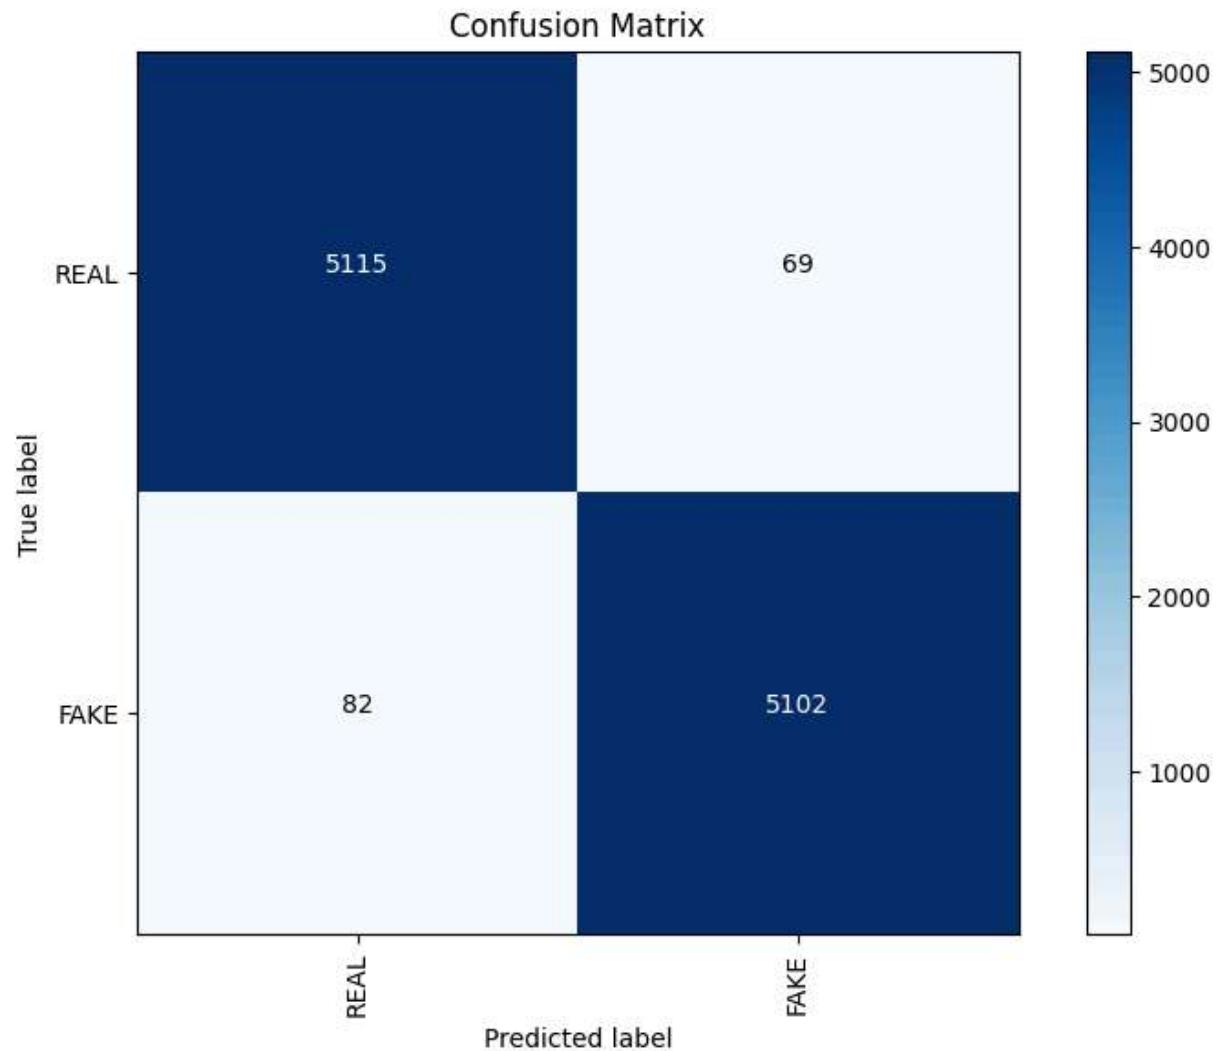
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()
    plt.show()

accuracy = accuracy_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred, average='macro')

print(f"Accuracy: {accuracy:.4f}")
print(f"F1 Score: {f1:.4f}")
if len(labels_list) <= 150:
    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, labels_list, figsize=(8, 6))
print()
print("Classification report:")
print()
print(classification_report(y_true, y_pred, target_names=labels_list, digits=4))
```

Accuracy: 0.9854

F1 Score: 0.9854



Classification report:

	precision	recall	f1-score	support
REAL	0.9842	0.9867	0.9855	5184
FAKE	0.9867	0.9842	0.9854	5184
accuracy			0.9854	10368
macro avg	0.9854	0.9854	0.9854	10368
weighted avg	0.9854	0.9854	0.9854	10368

In [37]:

```
trainer.save_model()
```

In [38]:

```
from transformers import pipeline  
  
pipe = pipeline('image-classification', model=model_name, device=0)
```

Using a slow image processor as `use\_fast` is unset and a slow processor was saved with this model. `use\_fast=True` will be the default behavior in v4.48, even if the model was saved with a slow processor. This will result in minor differences in outputs. You'll still be able to use a slow processor with `use\_fast=False`.

Device set to use cuda:0

In [39]:

```
image = test_data[1]["image"]  
  
. . .  
image
```

Out[39]:



In [40]:

```
pipe(image)
```

Out[40]:

```
[{'label': 'REAL', 'score': 0.9987921118736267},  
 {'label': 'FAKE', 'score': 0.0012079098960384727}]
```

In [41]:

```
id2label[test_data[1]["label"]]
```

Out[41]:

```
'REAL'
```

In [ ]:

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load
```

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking "Run" or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input/clothes-dataset'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside this notebook
```

```
07d7154538.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/kemeja_dress_putih_oversize_1697091330_37c86162_progressive_thumbnail.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/e59186b6-abfa-4f6a-9e7b-2c7e2cf3fd22.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/3193786c-22f1-4ea9-9e84-8ee47c138a8d.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/dress_pita_depan_1697103046_c7cc20e1_progressive_thumbnail.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/hm_dress_1697109054_8b47e0cb_progressive_thumbnail.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/fc5051e5-59aa-44e8-88dc-151981e61dcc.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/dress_cantik_bahan_jersey_prem_1697106977_53491f82_progressive_thumbnail.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/love_bonito_womans_midi_dress_1697087445_309021c7_progressive_thumbnail.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/f68947a8-6f6b-4461-ab13-81c285908509.jpg
/kaggle/input/clothes-dataset/Clothes_Dataset/Gaun/dress_ungu_nanipa_78_1697
```

```
In [2]: !pip install evaluate
import evaluate # for evaluating functions
print(evaluate.__version__)
```

```
Collecting evaluate
  Downloading evaluate-0.4.3-py3-none-any.whl.metadata (9.2 kB)
Requirement already satisfied: datasets>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (3.2.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.26.4)
Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.3.8)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.2.3)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.32.3)
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from evaluate) (4.67.1)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from evaluate) (3.5.0)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.70.16)
Requirement already satisfied: fsspec>=2021.05.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2021.05.0)
```

```
In [3]: import gc # garbage collection
import numpy as np
import pandas as pd
import itertools # for iterators and looping
from collections import Counter
```

```
In [4]: import matplotlib.pyplot as plt
from sklearn.metrics import (accuracy_score,
                             roc_auc_score,
                             confusion_matrix,
                             classification_report,
                             f1_score)
```

```
In [5]: from imblearn.over_sampling import RandomOverSampler
import accelerate
import evaluate
from transformers import (TrainingArguments,
                           Trainer,
                           ViTImageProcessor,
                           ViTForImageClassification,
                           DefaultDataCollator
)
```

```
In [7]: from PIL import ImageFile  
ImageFile.LOAD_TRUNCATED_IMAGES = True
```

```
In [8]: image_dict = {}

from pathlib import Path
from tqdm import tqdm
import os

file_names = []
labels = []
```

```
In [9]: label_dict = {
    "Blazer": "Blazer",
    "Celana_Panjang": "Long Pants",
    "Celana_Pendek": "Shorts",
    "Gaun": "Dresses",
    "Hoodie": "Hoodie",
    "Jaket": "Jacket",
    "Jaket_Denim": "Denim Jacket",
    "Jaket_Olahraga": "Sports Jacket",
    "Jeans": "Jeans",
    "Kaos": "T-shirt",
    "Kemeja": "Shirt",
    "Mantel": "Coat",
    "Polo": "Polo",
    "Rok": "Skirt",
    "Sweter": "Sweater",
}
```

```
In [10]: for file in tqdm(sorted((Path('/kaggle/input/clothes-dataset/Clothes_Dataset/').glob('*.jpg')))):
    label = str(file).split('/')[-2] # Extract the label from the file path
    labels.append(label_dict[label]) # Add the label to the list
    file_names.append(str(file)) # Add the file path to the list

print(len(file_names), len(labels))

df = pd.DataFrame.from_dict({"image": file_names, "label": labels})
print(df.shape)
```

100% |██████████| 7500/7500 [00:00<00:00, 416757.59it/s]

7500 7500  
(7500, 2)

In [11]: `df.head()`

Out[11]:

	image	label
0	/kaggle/input/clothes-dataset/Clothes_Dataset/...	Blazer
1	/kaggle/input/clothes-dataset/Clothes_Dataset/...	Blazer
2	/kaggle/input/clothes-dataset/Clothes_Dataset/...	Blazer
3	/kaggle/input/clothes-dataset/Clothes_Dataset/...	Blazer
4	/kaggle/input/clothes-dataset/Clothes_Dataset/...	Blazer

In [12]: `df["label"].value_counts()`

Out[12]:

label	count
Blazer	500
Long Pants	500
Shorts	500
Dresses	500
Hoodie	500
Jacket	500
Denim Jacket	500
Sports Jacket	500
Jeans	500
T-shirt	500
Shirt	500
Coat	500
Polo	500
Skirt	500
Sweater	500

Name: count, dtype: int64

In [13]:

```
from datasets import Dataset
from datasets import Image
dataset = Dataset.from_pandas(df).cast_column("image",Image())
```

```
In [14]: dataset[50]["image"]
```

Out[14]:





```
In [15]: labels_subset = labels[:5]
print(labels_subset)

['Blazer', 'Blazer', 'Blazer', 'Blazer', 'Blazer']
```

```
In [16]: labels_list = sorted(list(set(labels)))
```

```
In [17]: label2id, id2label = dict(),dict()

for i,label in enumerate(labels_list):
    label2id[label] = i
    id2label[i] = label

print(id2label)
print(label2id)
```

```
{0: 'Blazer', 1: 'Coat', 2: 'Denim Jacket', 3: 'Dresses', 4: 'Hoodie', 5: 'Jack
et', 6: 'Jeans', 7: 'Long Pants', 8: 'Polo', 9: 'Shirt', 10: 'Shorts', 11: 'Ski
rt', 12: 'Sports Jacket', 13: 'Sweater', 14: 'T-shirt'}
{'Blazer': 0, 'Coat': 1, 'Denim Jacket': 2, 'Dresses': 3, 'Hoodie': 4, 'Jacket': 5, 'Jeans': 6, 'Long Pants': 7, 'Polo': 8, 'Shirt': 9, 'Shorts': 10, 'Skirt': 11, 'Sports Jacket': 12, 'Sweater': 13, 'T-shirt': 14}
```

```
In [18]: import datasets
from datasets import ClassLabel
ClassLabels = ClassLabel(num_classes = len(labels_list),names = labels_list)
def map_label2id(example):
    example['label'] = ClassLabels.str2int(example['label'])
    return example
```

```
In [19]: dataset = dataset.map(map_label2id, batched=True)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
In [20]: dataset = dataset.cast_column('label', ClassLabels)

dataset = dataset.train_test_split(test_size=0.4, shuffle=True, stratify_by_column='label')

train_data = dataset['train']

test_data = dataset['test']
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
In [21]: model_str = 'google/vit-base-patch16-224-in21k'
```

```
In [22]: processor = ViTImageProcessor.from_pretrained(model_str)
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
In [23]: image_mean, image_std = processor.image_mean, processor.image_std
```

```
In [24]: size = processor.size["height"]
print("Size: ", size)
```

Size: 224

```
In [25]: normalize = Normalize(mean=image_mean, std=image_std)
```

```
In [26]: _train_transformers = Compose([
    Resize((size, size)),
    RandomRotation(90),
    RandomAdjustSharpness(2),
    RandomHorizontalFlip(0.5),
    ToTensor(),
    normalize
])
```

```
In [27]: _val_transforms = Compose(
    [
        Resize((size, size)),
        ToTensor(),
        normalize
    ]
)
```

```
In [28]: def train_transforms(examples):
    examples['pixel_values'] = [_train_transforms(image.convert("RGB")) for image in examples]
    return examples
```

```
In [29]: def val_transforms(examples):
    examples['pixel_values'] = [_val_transforms(image.convert("RGB")) for image in examples]
    return examples
```

```
In [30]: train_data.set_transform(train_transforms)
```

```
In [31]: test_data.set_transform(val_transforms)
```

```
In [32]: def collate_fn(examples):
    pixel_values = torch.stack([example["pixel_values"] for example in examples])
    labels = torch.tensor([example['label'] for example in examples])
    return {"pixel_values": pixel_values, "labels": labels}
```

```
In [33]: model = ViTForImageClassification.from_pretrained(model_str, num_labels=len(label2id))
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

Some weights of ViTForImageClassification were not initialized from the model checkpoint at google/vit-base-patch16-224-in21k and are newly initialized: ['classifier.bias', 'classifier.weight']  
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
In [34]: model.config.id2label = id2label
model.config.label2id = label2id
```

```
In [35]: print(model.num_parameters(only_trainable=True) / 1e6)
```

85.810191

```
In [36]: accuracy = evaluate.load("accuracy")
def compute_metrics(eval_pred):
    predictions = eval_pred.predictions
    label_ids = eval_pred.label_ids
    predicted_labels = predictions.argmax(axis=1)
    acc_score = accuracy.compute(predictions=predicted_labels, references=label_ids)
    return {
        "accuracy": acc_score
    }
```

A Jupyter widget could not be displayed because the widget state could not be found. This could happen if the kernel storing the widget is no longer available, or if the widget state was not saved in the notebook. You may be able to create the widget by running the appropriate cells.

```
In [37]: metric_name = "accuracy"
model_name = "clothes_image_detection"
num_train_epochs = 10
args = TrainingArguments(
    output_dir=model_name,
    logging_dir='./logs',
    evaluation_strategy="epoch",
    learning_rate=3e-6,
    per_device_train_batch_size=32,
    per_device_eval_batch_size=8,
    num_train_epochs=num_train_epochs,
    weight_decay=0.02,# it is use to prevent overfitting
    warmup_steps=50,
    remove_unused_columns=False,
    save_strategy='epoch',
    load_best_model_at_end=True,
    save_total_limit=1,
    report_to="none"
)
```

```
/usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1575: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.46 of 🤗 Transformers. Use `eval_strategy` instead
  warnings.warn(
```

```
In [38]: trainer = Trainer(
    model,
    args,
    train_dataset=train_data,
    eval_dataset=test_data,
    data_collator=collate_fn,
    compute_metrics=compute_metrics,
    tokenizer=processor,
)
```

```
<ipython-input-38-9c13e1b762e1>:1: FutureWarning: `tokenizer` is deprecated and
will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class`
instead.
  trainer = Trainer(
```

In [39]: `trainer.evaluate()`

```
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.  
    warnings.warn(
```

[188/188 05:08]

Out[39]: `{'eval_loss': 2.714015245437622,  
 'eval_accuracy': 0.07233333333333333,  
 'eval_runtime': 76.9229,  
 'eval_samples_per_second': 39.0,  
 'eval_steps_per_second': 2.444}`

In [40]: `trainer.train()`

```
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.  
    warnings.warn(
```

[710/710 35:46, Epoch 10/10]

Epoch	Training Loss	Validation Loss	Accuracy
1	No log	2.688771	0.112667
2	No log	2.649584	0.200333
3	No log	2.607557	0.311000
4	No log	2.563295	0.397667
5	No log	2.523514	0.456667
6	No log	2.490482	0.509667
7	No log	2.464117	0.543333
8	2.574500	2.445119	0.562333
9	2.574500	2.433961	0.577667
10	2.574500	2.430238	0.584000

```
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
```

**Out[40]:** TrainOutput(global\_step=710, training\_loss=2.5268360406580106, metrics={'train\_runtime': 2150.1722, 'train\_samples\_per\_second': 20.929, 'train\_steps\_per\_second': 0.33, 'total\_flos': 3.48754583632896e+18, 'train\_loss': 2.5268360406580106, 'epoch': 10.0})

In [41]: `trainer.evaluate()`

```
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
```

Out[41]: `{'eval_loss': 2.4302384853363037,
 'eval_accuracy': 0.584,
 'eval_runtime': 61.3954,
 'eval_samples_per_second': 48.864,
 'eval_steps_per_second': 3.062,
 'epoch': 10.0}`

In [42]: `outputs = trainer.predict(test_data)
print(outputs.metrics)`

```
/usr/local/lib/python3.10/dist-packages/torch/nn/parallel/_functions.py:71: Use
rWarning: Was asked to gather along dimension 0, but all input tensors were sca
lars; will instead unsqueeze and return a vector.
    warnings.warn(
{'test_loss': 2.4302384853363037, 'test_accuracy': 0.584, 'test_runtime': 61.60
22, 'test_samples_per_second': 48.7, 'test_steps_per_second': 3.052}
```

```
In [43]: y_true = outputs.label_ids
y_pred = outputs.predictions.argmax(1)
def plot_confusion_matrix(cm, classes, title='Confusion Matrix', cmap=plt.cm.Blue
    """
        This function plots a confusion matrix.

    Parameters:
        cm (array-like): Confusion matrix as returned by sklearn.metrics.confusion_matrix.
        classes (list): List of class names, e.g., ['Class 0', 'Class 1'].
        title (str): Title for the plot.
        cmap (matplotlib colormap): Colormap for the plot.
    """
    plt.figure(figsize=figsize)
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)
    fmt = '.0f'
    thresh = cm.max() / 2.0
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt), horizontalalignment="center", color="white" if cm[i, j] > thresh else "black")
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()
    plt.show()

accuracy = accuracy_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred, average='macro')

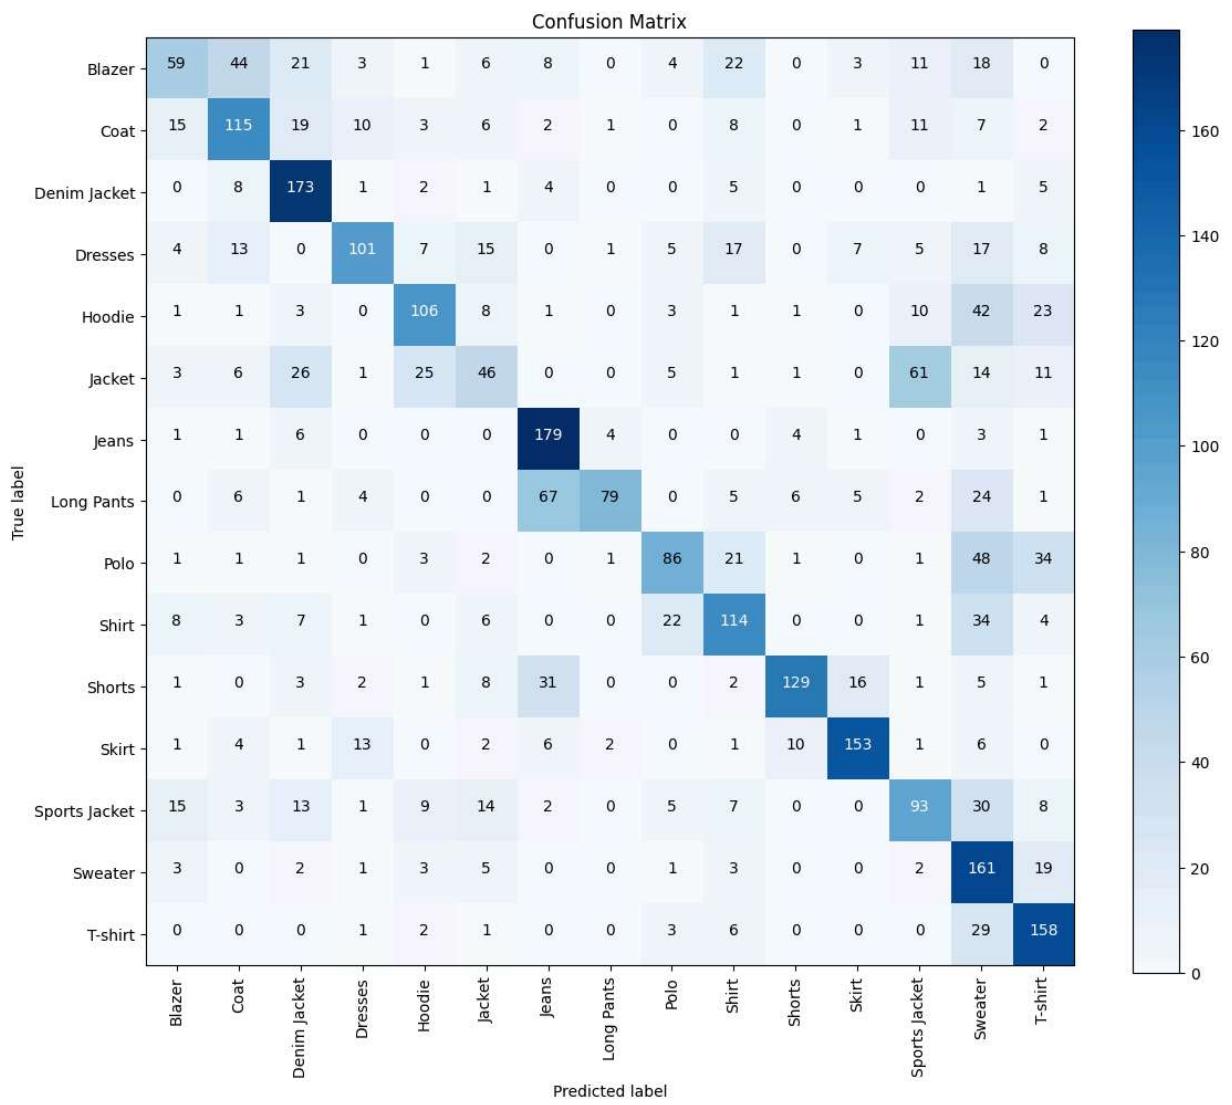
print(f"Accuracy: {accuracy:.4f}")
print(f"F1 Score: {f1:.4f}")

if len(labels_list) <= 150:
    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, labels_list, figsize=(12, 10))

print()
print("Classification report:")
print()
print(classification_report(y_true, y_pred, target_names=labels_list, digits=4))
```

Accuracy: 0.5840

F1 Score: 0.5756



Classification report:

	precision	recall	f1-score	support
Blazer	0.5268	0.2950	0.3782	200
Coat	0.5610	0.5750	0.5679	200
Denim Jacket	0.6268	0.8650	0.7269	200
Dresses	0.7266	0.5050	0.5959	200
Hoodie	0.6543	0.5300	0.5856	200
Jacket	0.3833	0.2300	0.2875	200
Jeans	0.5967	0.8950	0.7160	200
Long Pants	0.8977	0.3950	0.5486	200
Polo	0.6418	0.4300	0.5150	200
Shirt	0.5352	0.5700	0.5521	200
Shorts	0.8487	0.6450	0.7330	200
Skirt	0.8226	0.7650	0.7927	200
Sports Jacket	0.4673	0.4650	0.4662	200
Sweater	0.3667	0.8050	0.5039	200
T-shirt	0.5745	0.7900	0.6653	200
accuracy			0.5840	3000
macro avg	0.6153	0.5840	0.5756	3000
weighted avg	0.6153	0.5840	0.5756	3000

In [44]: `trainer.save_model()`

In [45]: `from transformers import pipeline  
pipe = pipeline('image-classification', model=model_name, device=0)`

Device set to use cuda:0

```
In [46]: image = test_data[1]["image"]
image
```

Out[46]:





In [47]: `pipe(image)`

Out[47]:

```
[{'label': 'Denim Jacket', 'score': 0.6254792809486389},  
 {'label': 'Shirt', 'score': 0.5571456551551819},  
 {'label': 'Sports Jacket', 'score': 0.5241682529449463},  
 {'label': 'Polo', 'score': 0.5233001112937927},  
 {'label': 'Jeans', 'score': 0.5214868187904358}]
```

In [48]: `id2label[test_data[1]["label"]]`

Out[48]: 'Denim Jacket'



In [1]: pip install seaborn

```
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.12.2)
Requirement already satisfied: numpy!=1.24.0,>=1.17 in /usr/local/lib/python3.10/dist-packages (from seaborn) (1.26.4)
Requirement already satisfied: pandas>=0.25 in /usr/local/lib/python3.10/dist-packages (from seaborn) (2.2.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in /usr/local/lib/python3.10/dist-packages (from seaborn) (3.7.5)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.55.3)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (24.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.9.0.post0)
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (1.3.8)
Requirement already satisfied: mkl_random in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (1.2.4)
Requirement already satisfied: mkl_umat in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (2025.0.1)
Requirement already satisfied: tbb4py in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (2022.0.0)
Requirement already satisfied: mkl-service in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (2.4.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.25->seaborn) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.25->seaborn) (2025.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.17.0)
Requirement already satisfied: intel-openmp>=2024 in /usr/local/lib/python3.10/dist-packages (from mkl->numpy!=1.24.0,>=1.17->seaborn) (2024.2.0)
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.10/dist-packages (from mkl->numpy!=1.24.0,>=1.17->seaborn) (2022.0.0)
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.10/dist-packages (from tbb==2022.*->mkl->numpy!=1.24.0,>=1.17->seaborn) (1.2.0)
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.10/dist-packages (from mkl_umat->numpy!=1.24.0,>=1.17->seaborn) (2024.2.0)
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.10/dist-packages (from intel-openmp>=2024->mkl->numpy!=1.24.0,>=1.17->seaborn) (2024.2.0)
```

Note: you may need to restart the kernel to use updated packages.

```
In [2]: pip install split-folders
```

```
Collecting split-folders
  Downloading split_folders-0.5.1-py3-none-any.whl.metadata (6.2 kB)
Downloaded split_folders-0.5.1-py3-none-any.whl (8.4 kB)
Installing collected packages: split-folders
Successfully installed split-folders-0.5.1
Note: you may need to restart the kernel to use updated packages.
```

```
In [3]: pip install numpy==1.23.4
```

```
Collecting numpy==1.23.4
  Downloading numpy-1.23.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_
  64.whl.metadata (2.3 kB)
  Downloading numpy-1.23.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_6
  4.whl (17.1 MB)
----- 17.1/17.1 MB 92.7 MB/s eta 0:00:00:
00:0100:01
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 1.26.4
    Uninstalling numpy-1.26.4:
      Successfully uninstalled numpy-1.26.4
ERROR: pip's dependency resolver does not currently take into account all the p
ackages that are installed. This behaviour is the source of the following depen
dency conflicts.
albucore 0.0.19 requires numpy>=1.24.4, but you have numpy 1.23.4 which is inco
mpatible.
albumentations 1.4.20 requires numpy>=1.24.4, but you have numpy 1.23.4 which i
s incompatible.
bayesian-optimization 2.0.3 requires numpy>=1.25, but you have numpy 1.23.4 whi
ch is incompatible.
bigframes 1.29.0 requires numpy>=1.24.0, but you have numpy 1.23.4 which is inc
ompatible.
chex 0.1.88 requires numpy>=1.24.1, but you have numpy 1.23.4 which is incompat
ible.
featuretools 1.31.0 requires numpy>=1.25.0, but you have numpy 1.23.4 which is incom
patible.
jax 0.4.33 requires numpy>=1.24, but you have numpy 1.23.4 which is incompatibl
e.
jaxlib 0.4.33 requires numpy>=1.24, but you have numpy 1.23.4 which is incompat
ible.
langchain 0.3.12 requires async-timeout<5.0.0,>=4.0.0; python_version < "3.11",
but you have async-timeout 5.0.1 which is incompatible.
mizani 0.13.1 requires numpy>=1.23.5, but you have numpy 1.23.4 which is incom
patible.
mkl-fft 1.3.8 requires numpy<1.27.0,>=1.26.4, but you have numpy 1.23.4 which i
s incompatible.
mkl-random 1.2.4 requires numpy<1.27.0,>=1.26.4, but you have numpy 1.23.4 whic
h is incompatible.
mkl-umath 0.1.1 requires numpy<1.27.0,>=1.26.4, but you have numpy 1.23.4 which
is incompatible.
mlxtend 0.23.3 requires scikit-learn>=1.3.1, but you have scikit-learn 1.2.2 wh
ich is incompatible.
pandas-gbq 0.25.0 requires google-api-core<3.0.0dev,>=2.10.2, but you have goog
le-api-core 1.34.1 which is incompatible.
pandas-stubs 2.2.2.240909 requires numpy>=1.23.5, but you have numpy 1.23.4 whi
ch is incompatible.
plotnine 0.14.4 requires matplotlib>=3.8.0, but you have matplotlib 3.7.5 which
is incompatible.
plotnine 0.14.4 requires numpy>=1.23.5, but you have numpy 1.23.4 which is inco
mpatible.
pyldavis 3.4.1 requires numpy>=1.24.2, but you have numpy 1.23.4 which is incom
patible.
pymc 5.19.1 requires numpy>=1.25.0, but you have numpy 1.23.4 which is incompat
ible.
scikit-image 0.25.0 requires numpy>=1.24, but you have numpy 1.23.4 which is in
compatible.
```

```
tensorflow 2.17.1 requires numpy<2.0.0,>=1.23.5; python_version <= "3.11", but
you have numpy 1.23.4 which is incompatible.
tensorflow-decision-forests 1.10.0 requires tensorflow==2.17.0, but you have te
nsorflow 2.17.1 which is incompatible.
woodwork 0.31.0 requires numpy>=1.25.0, but you have numpy 1.23.4 which is inco
mpatible.
xarray 2024.11.0 requires numpy>=1.24, but you have numpy 1.23.4 which is incom
patible.
Successfully installed numpy-1.23.4
Note: you may need to restart the kernel to use updated packages.
```

```
In [4]: import numpy as np
import os
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import backend as K
from tensorflow.keras import regularizers
from tensorflow.keras.layers import *
from tensorflow.keras.callbacks import *
from tensorflow.keras.optimizers import *
from tensorflow.keras.models import load_model, Model
from tensorflow.python.keras.utils import conv_utils
from tensorflow.keras import regularizers, constraints, initializers
import pandas as pd
from tensorflow.keras.layers import GlobalAveragePooling2D
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report

tf.random.set_seed(42)
np.random.seed(42)
```

```
In [5]: train_path = ("/kaggle/input/brain-tumor-mri-dataset/Training")
val_path = ("/kaggle/input/brain-tumor-mri-dataset/Testing")
```

```
In [6]: batch_size = 32
img_height = 128
img_width = 128
no_of_clases = 4
clses_name = ['glioma', 'meningioma', 'notumor', 'pituitary']
input_shape = (img_height, img_width, 3)
```

```
In [7]: train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=20,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest'  
)  
  
val_datagen = ImageDataGenerator(rescale=1./255)  
  
train_generator = train_datagen.flow_from_directory(  
    train_path,  
    target_size=(img_height, img_width),  
    batch_size=batch_size,  
    shuffle=True,  
    class_mode='categorical'  
)  
  
validation_generator = val_datagen.flow_from_directory(  
    val_path,  
    target_size=(img_height, img_width),  
    batch_size=batch_size,  
    shuffle=False, # No need to shuffle validation data  
    class_mode='categorical'  
)
```

Found 5712 images belonging to 4 classes.  
Found 1311 images belonging to 4 classes.

```
In [8]: print('train classes')  
print(train_generator.class_indices)  
print('validation classes')  
print(validation_generator.class_indices)
```

```
train classes  
{'glioma': 0, 'meningioma': 1, 'notumor': 2, 'pituitary': 3}  
validation classes  
{'glioma': 0, 'meningioma': 1, 'notumor': 2, 'pituitary': 3}
```

```
In [9]: model = tf.keras.models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=input_shape),
    MaxPooling2D(2, 2),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(no_of_clases, activation='softmax') # Output Layer
])
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [10]: model.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	P
conv2d (Conv2D)	(None, 126, 126, 32)	
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	
conv2d_1 (Conv2D)	(None, 61, 61, 64)	
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	
conv2d_2 (Conv2D)	(None, 28, 28, 128)	
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	
flatten (Flatten)	(None, 25088)	
dense (Dense)	(None, 128)	3,2
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 4)	



Total params: 3,305,156 (12.61 MB)

Trainable params: 3,305,156 (12.61 MB)

Non-trainable params: 0 (0.00 B)

```
In [11]: epochs = 10
history = model.fit(
    train_generator,
    steps_per_epoch=len(train_generator),
    epochs=epochs,
    validation_data=validation_generator,
    validation_steps=len(validation_generator)
)
```

Epoch 1/10

```
/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dat
aset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__
_init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_m
ultiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as t
hey will be ignored.
    self._warn_if_super_not_called()
```

```
179/179 ━━━━━━━━ 58s 289ms/step - accuracy: 0.4549 - loss: 1.1587 -
val_accuracy: 0.4630 - val_loss: 2.0576
```

Epoch 2/10

```
179/179 ━━━━━━━━ 0s 93us/step - accuracy: 0.0000e+00 - loss: 0.0000
e+00
```

Epoch 3/10

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data;
interrupting training. Make sure that your dataset or generator can generate at
least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` f
unction when building your dataset.
```

```
    self.gen.throw(typ, value, traceback)
```

```
179/179 ━━━━━━━━ 32s 171ms/step - accuracy: 0.6867 - loss: 0.7860 -
val_accuracy: 0.7162 - val_loss: 0.6928
```

Epoch 4/10

```
179/179 ━━━━━━━━ 0s 63us/step - accuracy: 0.0000e+00 - loss: 0.0000
e+00
```

Epoch 5/10

```
179/179 ━━━━━━━━ 32s 175ms/step - accuracy: 0.7363 - loss: 0.6905 -
val_accuracy: 0.5858 - val_loss: 1.2726
```

Epoch 6/10

```
179/179 ━━━━━━━━ 0s 62us/step - accuracy: 0.0000e+00 - loss: 0.0000
e+00
```

Epoch 7/10

```
179/179 ━━━━━━━━ 32s 170ms/step - accuracy: 0.7663 - loss: 0.6190 -
val_accuracy: 0.7811 - val_loss: 0.5902
```

Epoch 8/10

```
179/179 ━━━━━━━━ 0s 67us/step - accuracy: 0.0000e+00 - loss: 0.0000
e+00
```

Epoch 9/10

```
179/179 ━━━━━━━━ 33s 180ms/step - accuracy: 0.7718 - loss: 0.6066 -
val_accuracy: 0.5820 - val_loss: 1.2488
```

Epoch 10/10

```
179/179 ━━━━━━━━ 0s 48us/step - accuracy: 0.0000e+00 - loss: 0.0000
e+00
```

```
In [12]: loss, accuracy = model.evaluate(validation_generator, steps=len(validation_generator))
print('Validation accuracy:', accuracy)
```

41/41 ━━━━━━━━━━ 3s 67ms/step - accuracy: 0.4610 - loss: 1.6857  
Validation accuracy: 0.5819984674453735

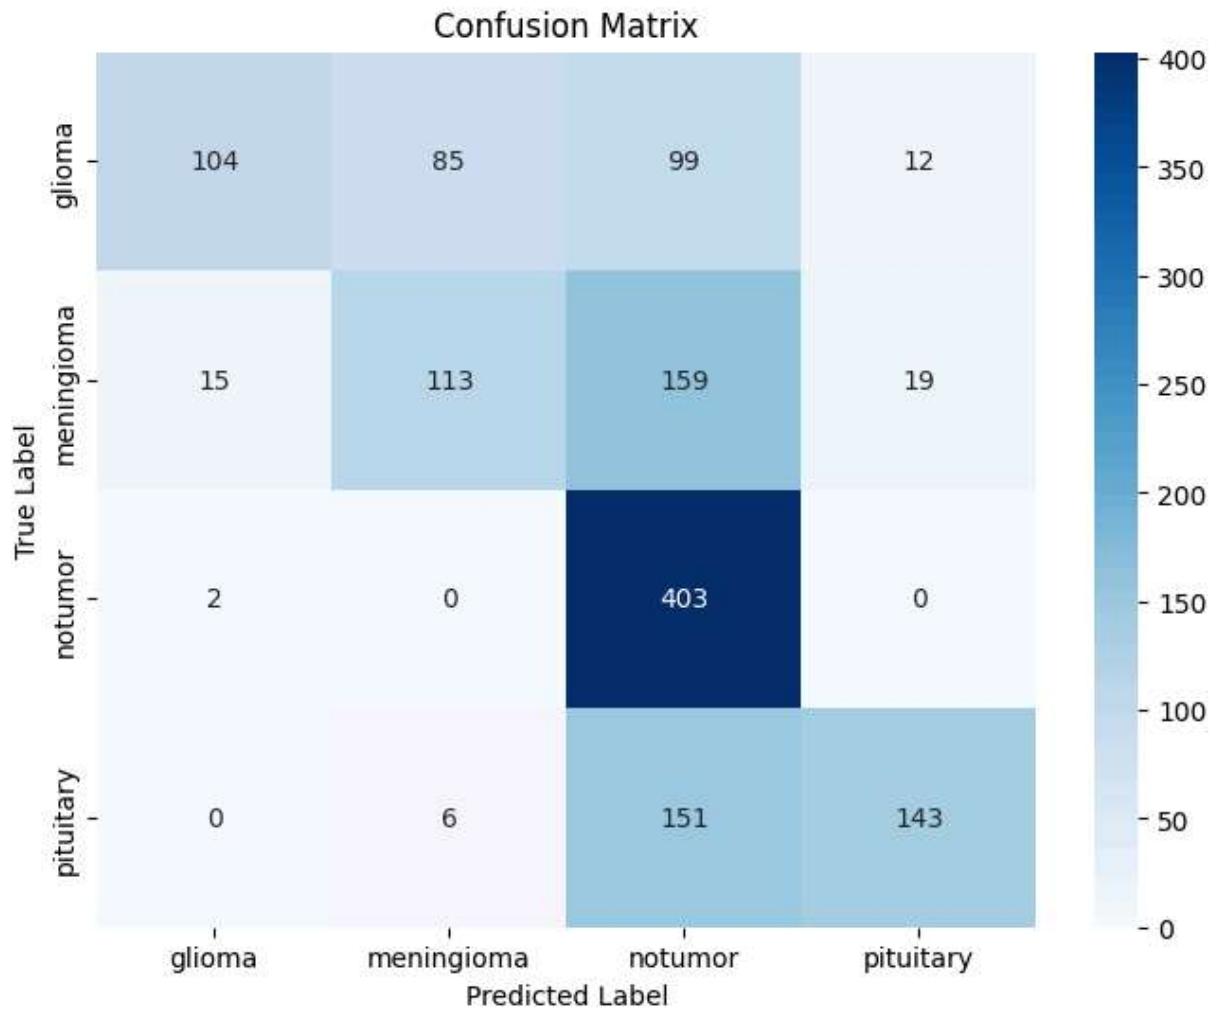
```
In [13]: # Classification Report and Confusion Matrix
y_pred = model.predict(validation_generator, steps=len(validation_generator))
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = validation_generator.classes
class_names = list(validation_generator.class_indices.keys())

print('\nClassification Report:\n', classification_report(y_true, y_pred_classes))
```

41/41 ━━━━━━━━━━ 3s 67ms/step

	precision	recall	f1-score	support
glioma	0.86	0.35	0.49	300
meningioma	0.55	0.37	0.44	306
notumor	0.50	1.00	0.66	405
pituitary	0.82	0.48	0.60	300
accuracy			0.58	1311
macro avg	0.68	0.55	0.55	1311
weighted avg	0.67	0.58	0.56	1311

```
In [14]: cm = confusion_matrix(y_true, y_pred_classes)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, ytick
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [ ]:
```



In [1]: pip install seaborn

```
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.12.2)
Requirement already satisfied: numpy!=1.24.0,>=1.17 in /usr/local/lib/python3.10/dist-packages (from seaborn) (1.26.4)
Requirement already satisfied: pandas>=0.25 in /usr/local/lib/python3.10/dist-packages (from seaborn) (2.2.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in /usr/local/lib/python3.10/dist-packages (from seaborn) (3.7.5)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.55.3)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (24.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.9.0.post0)
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (1.3.8)
Requirement already satisfied: mkl_random in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (1.2.4)
Requirement already satisfied: mkl_umat in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (2025.0.1)
Requirement already satisfied: tbb4py in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (2022.0.0)
Requirement already satisfied: mkl-service in /usr/local/lib/python3.10/dist-packages (from numpy!=1.24.0,>=1.17->seaborn) (2.4.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.25->seaborn) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.25->seaborn) (2025.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.17.0)
Requirement already satisfied: intel-openmp>=2024 in /usr/local/lib/python3.10/dist-packages (from mkl->numpy!=1.24.0,>=1.17->seaborn) (2024.2.0)
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.10/dist-packages (from mkl->numpy!=1.24.0,>=1.17->seaborn) (2022.0.0)
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.10/dist-packages (from tbb==2022.*->mkl->numpy!=1.24.0,>=1.17->seaborn) (1.2.0)
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.10/dist-packages (from mkl_umat->numpy!=1.24.0,>=1.17->seaborn) (2024.2.0)
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.10/dist-packages (from intel-openmp>=2024->mkl->numpy!=1.24.0,>=1.17->seaborn) (2024.2.0)
```

Note: you may need to restart the kernel to use updated packages.

```
In [2]: pip install numpy==1.23.4
```

```
Collecting numpy==1.23.4
  Downloading numpy-1.23.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_
  64.whl.metadata (2.3 kB)
  Downloading numpy-1.23.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_6
  4.whl (17.1 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 17.1/17.1 MB 75.8 MB/s eta 0:00:00:
00:0100:01
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 1.26.4
    Uninstalling numpy-1.26.4:
      Successfully uninstalled numpy-1.26.4
ERROR: pip's dependency resolver does not currently take into account all the p
ackages that are installed. This behaviour is the source of the following depen
dency conflicts.
albucore 0.0.19 requires numpy>=1.24.4, but you have numpy 1.23.4 which is inco
mpatible.
albumentations 1.4.20 requires numpy>=1.24.4, but you have numpy 1.23.4 which i
s incompatible.
bayesian-optimization 2.0.3 requires numpy>=1.25, but you have numpy 1.23.4 whi
ch is incompatible.
bigframes 1.29.0 requires numpy>=1.24.0, but you have numpy 1.23.4 which is inc
ompatible.
chex 0.1.88 requires numpy>=1.24.1, but you have numpy 1.23.4 which is incompat
ible.
featuretools 1.31.0 requires numpy>=1.25.0, but you have numpy 1.23.4 which is incom
patible.
jax 0.4.33 requires numpy>=1.24, but you have numpy 1.23.4 which is incompatibl
e.
jaxlib 0.4.33 requires numpy>=1.24, but you have numpy 1.23.4 which is incompat
ible.
langchain 0.3.12 requires async-timeout<5.0.0,>=4.0.0; python_version < "3.11",
but you have async-timeout 5.0.1 which is incompatible.
mizani 0.13.1 requires numpy>=1.23.5, but you have numpy 1.23.4 which is incomp
atible.
mkl-fft 1.3.8 requires numpy<1.27.0,>=1.26.4, but you have numpy 1.23.4 which i
s incompatible.
mkl-random 1.2.4 requires numpy<1.27.0,>=1.26.4, but you have numpy 1.23.4 whic
h is incompatible.
mkl-umath 0.1.1 requires numpy<1.27.0,>=1.26.4, but you have numpy 1.23.4 which
is incompatible.
mlxtend 0.23.3 requires scikit-learn>=1.3.1, but you have scikit-learn 1.2.2 wh
ich is incompatible.
pandas-gbq 0.25.0 requires google-api-core<3.0.0dev,>=2.10.2, but you have goog
le-api-core 1.34.1 which is incompatible.
pandas-stubs 2.2.2.240909 requires numpy>=1.23.5, but you have numpy 1.23.4 whi
ch is incompatible.
plotnine 0.14.4 requires matplotlib>=3.8.0, but you have matplotlib 3.7.5 which
is incompatible.
plotnine 0.14.4 requires numpy>=1.23.5, but you have numpy 1.23.4 which is inco
mpatible.
pyldavis 3.4.1 requires numpy>=1.24.2, but you have numpy 1.23.4 which is incom
patible.
pymc 5.19.1 requires numpy>=1.25.0, but you have numpy 1.23.4 which is incompat
ible.
scikit-image 0.25.0 requires numpy>=1.24, but you have numpy 1.23.4 which is in
compatible.
```

```
tensorflow 2.17.1 requires numpy<2.0.0,>=1.23.5; python_version <= "3.11", but
you have numpy 1.23.4 which is incompatible.
tensorflow-decision-forests 1.10.0 requires tensorflow==2.17.0, but you have te
nsorflow 2.17.1 which is incompatible.
woodwork 0.31.0 requires numpy>=1.25.0, but you have numpy 1.23.4 which is inco
mpatible.
xarray 2024.11.0 requires numpy>=1.24, but you have numpy 1.23.4 which is incom
patible.
Successfully installed numpy-1.23.4
Note: you may need to restart the kernel to use updated packages.
```

```
In [3]: import numpy as np
import os
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
print(tf.__version__)
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import backend as C
from tensorflow.keras import regularizers
from tensorflow.keras.layers import *
from tensorflow.keras.callbacks import *
from tensorflow.keras.optimizers import *
from tensorflow.keras.models import load_model,Model
from tensorflow.python.keras.utils import conv_utils
from tensorflow.keras import regularizers, constraints, initializers
import pandas as pd
from tensorflow.keras.layers import GlobalAveragePooling2D
# import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

2.17.1

```
In [4]: train_path="/kaggle/input/cards-image-datasetclassification/train"
val_path="/kaggle/input/cards-image-datasetclassification/valid"
test_path="/kaggle/input/cards-image-datasetclassification/test"
batch_size = 32
img_height = 640
img_width = 640
no_of_clases = 53
classes_name = [
    'AS', '2S', '3S', '4S', '5S', '6S', '7S', '8S', '9S', '10S',
    'JS', 'QS', 'KS', # Spades
    'AH', '2H', '3H', '4H', '5H', '6H', '7H', '8H', '9H', '10H',
    'JH', 'QH', 'KH', # Hearts
    'AD', '2D', '3D', '4D', '5D', '6D', '7D', '8D', '9D', '10D',
    'JD', 'QD', 'KD', # Diamonds
    'AC', '2C', '3C', '4C', '5C', '6C', '7C', '8C', '9C', '10C',
    'JC', 'QC', 'KC', # Clubs
    'BACK'           # Card Back
]

input_shape = (img_height,img_width,1)

datagen = ImageDataGenerator(rescale = 1./255,featurewise_center=True,horizontal_
train_generator = datagen.flow_from_directory(train_path,target_size=(img_height,
validation_generator = datagen.flow_from_directory(val_path,target_size=(img_heig
print('train classes')
print(train_generator.class_indices)
print('train labels')
print(train_generator.labels)
print('validation claes')
print(validation_generator.class_indices)
print('validation labels')
print(validation_generator.labels)
```

```
Found 7624 images belonging to 53 classes.
```

```
Found 265 images belonging to 53 classes.
```

```
train classes
```

```
{'ace of clubs': 0, 'ace of diamonds': 1, 'ace of hearts': 2, 'ace of spades': 3, 'eight of clubs': 4, 'eight of diamonds': 5, 'eight of hearts': 6, 'eight of spades': 7, 'five of clubs': 8, 'five of diamonds': 9, 'five of hearts': 10, 'five of spades': 11, 'four of clubs': 12, 'four of diamonds': 13, 'four of hearts': 14, 'four of spades': 15, 'jack of clubs': 16, 'jack of diamonds': 17, 'jack of hearts': 18, 'jack of spades': 19, 'joker': 20, 'king of clubs': 21, 'king of diamonds': 22, 'king of hearts': 23, 'king of spades': 24, 'nine of clubs': 25, 'nine of diamonds': 26, 'nine of hearts': 27, 'nine of spades': 28, 'queen of clubs': 29, 'queen of diamonds': 30, 'queen of hearts': 31, 'queen of spades': 32, 'seven of clubs': 33, 'seven of diamonds': 34, 'seven of hearts': 35, 'seven of spades': 36, 'six of clubs': 37, 'six of diamonds': 38, 'six of hearts': 39, 'six of spades': 40, 'ten of clubs': 41, 'ten of diamonds': 42, 'ten of hearts': 43, 'ten of spades': 44, 'three of clubs': 45, 'three of diamonds': 46, 'three of hearts': 47, 'three of spades': 48, 'two of clubs': 49, 'two of diamonds': 50, 'two of hearts': 51, 'two of spades': 52}
```

```
train labels
```

```
[ 0  0  0 ... 52 52 52]
```

```
validation classes
```

```
{'ace of clubs': 0, 'ace of diamonds': 1, 'ace of hearts': 2, 'ace of spades': 3, 'eight of clubs': 4, 'eight of diamonds': 5, 'eight of hearts': 6, 'eight of spades': 7, 'five of clubs': 8, 'five of diamonds': 9, 'five of hearts': 10, 'five of spades': 11, 'four of clubs': 12, 'four of diamonds': 13, 'four of hearts': 14, 'four of spades': 15, 'jack of clubs': 16, 'jack of diamonds': 17, 'jack of hearts': 18, 'jack of spades': 19, 'joker': 20, 'king of clubs': 21, 'king of diamonds': 22, 'king of hearts': 23, 'king of spades': 24, 'nine of clubs': 25, 'nine of diamonds': 26, 'nine of hearts': 27, 'nine of spades': 28, 'queen of clubs': 29, 'queen of diamonds': 30, 'queen of hearts': 31, 'queen of spades': 32, 'seven of clubs': 33, 'seven of diamonds': 34, 'seven of hearts': 35, 'seven of spades': 36, 'six of clubs': 37, 'six of diamonds': 38, 'six of hearts': 39, 'six of spades': 40, 'ten of clubs': 41, 'ten of diamonds': 42, 'ten of hearts': 43, 'ten of spades': 44, 'three of clubs': 45, 'three of diamonds': 46, 'three of hearts': 47, 'three of spades': 48, 'two of clubs': 49, 'two of diamonds': 50, 'two of hearts': 51, 'two of spades': 52}
```

```
validation labels
```

```
[ 0  0  0  0  0  1  1  1  1  1  2  2  2  2  2  3  3  3  3  3  3  4  4  4  4
 4  5  5  5  5  6  6  6  6  6  7  7  7  7  7  7  8  8  8  8  8  8  9  9  9
 9  9 10 10 10 10 11 11 11 11 11 12 12 12 12 12 13 13 13 13 13 13 14 14
14 14 14 15 15 15 15 15 16 16 16 16 16 16 17 17 17 17 17 17 18 18 18 18 18
19 19 19 19 20 20 20 20 21 21 21 21 21 22 22 22 22 22 23 23 23 23 23
24 24 24 24 24 25 25 25 25 25 26 26 26 26 26 27 27 27 27 27 28 28 28 28
28 29 29 29 29 30 30 30 30 30 31 31 31 31 31 32 32 32 32 32 33 33 33
33 33 34 34 34 34 35 35 35 35 36 36 36 36 36 37 37 37 37 37 38 38 38
38 38 38 39 39 39 39 39 40 40 40 40 40 41 41 41 41 41 42 42 42 42 42 43
43 43 43 43 44 44 44 44 44 45 45 45 45 45 46 46 46 46 46 47 47 47 47 47
48 48 48 48 48 49 49 49 49 49 50 50 50 50 50 51 51 51 51 51 52 52 52 52
52]
```

```
In [5]: pretrained_model = tf.keras.applications.ResNet152V2(
    input_shape=(640, 640, 3),
    include_top=False,
    weights='imagenet',
    pooling='avg')

pretrained_model.trainable = False
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5) ([https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5))

234545216/234545216 ————— 1s 0us/step

```
In [6]: inputs = pretrained_model.input
outputs1 = tf.keras.layers.Dense(1024, activation='relu')(pretrained_model.output)
dropout = tf.keras.layers.Dropout(0.5)(outputs1)
output2 = tf.keras.layers.Dense(1024, activation = 'relu')(dropout)
dropout2 = tf.keras.layers.Dropout(0.5)(output2)
outputs = tf.keras.layers.Dense(53, activation = 'softmax')(dropout2)
model = tf.keras.Model(inputs, outputs)
print(model.summary())
```

Model: "functional"

Layer (type)	Output Shape	Param #	Connec
input_layer (InputLayer)	(None, 640, 640, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 646, 646, 3)	0	input_
conv1_conv (Conv2D)	(None, 320, 320, 64)	9,472	conv1_
pool1_pad (ZeroPadding2D)	(None, 322, 322, 64)	0	conv1_
pool1_pool (MaxPooling2D)	(None, 160, 160, 64)	0	pool1_
conv2_block1_preact_bn (BatchNormalization)	(None, 160, 160, 64)	256	pool1_
conv2_block1_preact_relu	(None, 160, 160, 64)	0	conv2_

```
In [7]: model.compile(optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, decay=1e-5),
```

/usr/local/lib/python3.10/dist-packages/keras/src/optimizers/base\_optimizer.py:33: UserWarning: Argument `decay` is no longer supported and will be ignored.  
warnings.warn(

```
In [8]: history = model.fit(train_generator, validation_data=validation_generator, epochs=10)

/usr/local/lib/python3.10/dist-packages/keras/src/legacy/preprocessing/image.py:1263: UserWarning: This ImageDataGenerator specifies `featurewise_center`, but it hasn't been fit on any training data. Fit it first by calling `.`.fit(numpy_data)`.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/keras/src/legacy/preprocessing/image.py:1273: UserWarning: This ImageDataGenerator specifies `featurewise_std_normalization`, but it hasn't been fit on any training data. Fit it first by calling `.`.fit(numpy_data)`.
    warnings.warn()

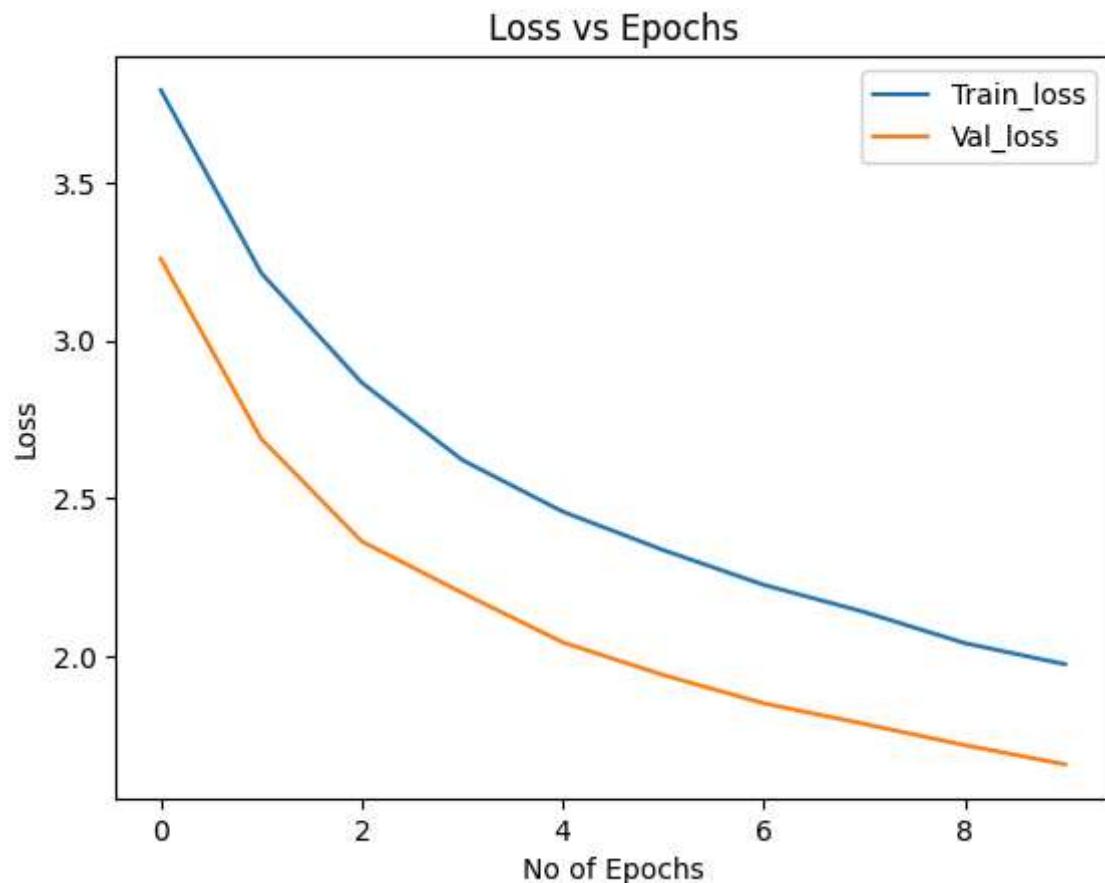
Epoch 1/10

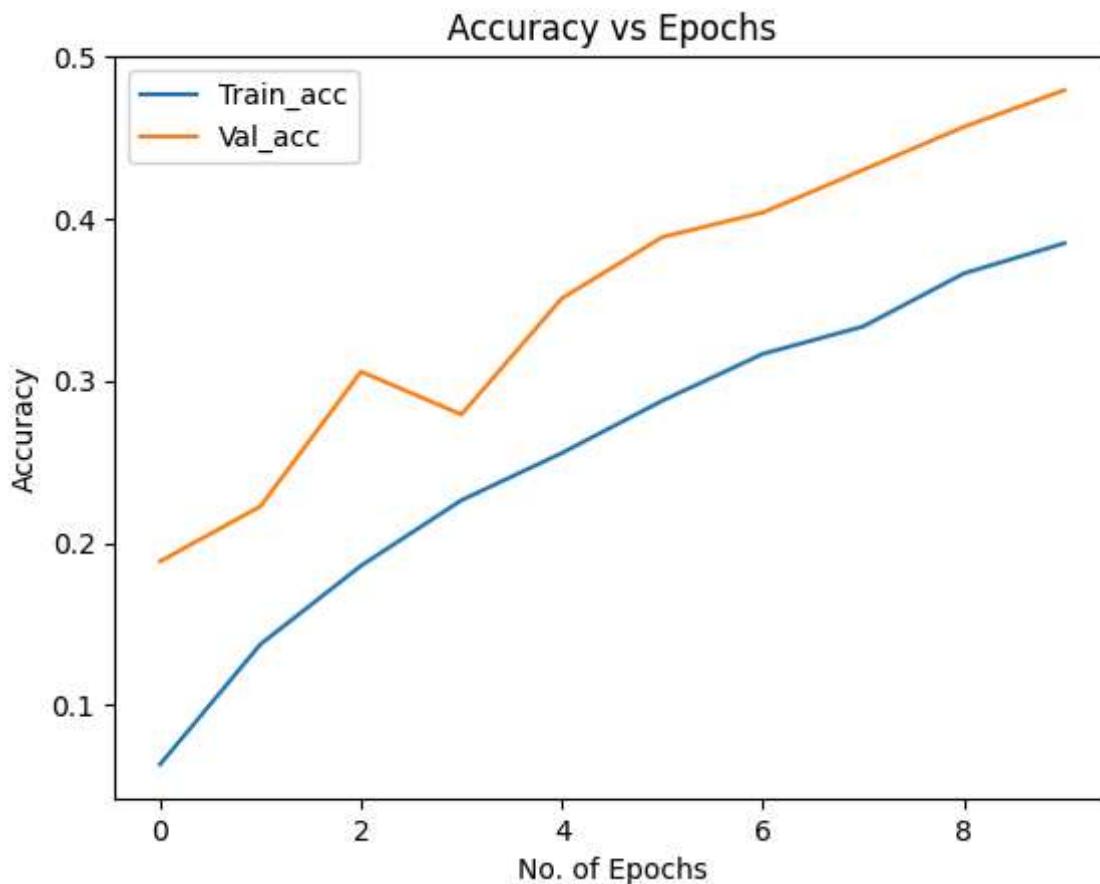
/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
    self._warn_if_super_not_called()

239/239 467s 2s/step - accuracy: 0.0396 - loss: 3.9823 - val_accuracy: 0.1887 - val_loss: 3.2589
Epoch 2/10
239/239 382s 2s/step - accuracy: 0.1206 - loss: 3.3266 - val_accuracy: 0.2226 - val_loss: 2.6875
Epoch 3/10
239/239 382s 2s/step - accuracy: 0.1797 - loss: 2.9279 - val_accuracy: 0.3057 - val_loss: 2.3631
Epoch 4/10
239/239 382s 2s/step - accuracy: 0.2289 - loss: 2.6432 - val_accuracy: 0.2792 - val_loss: 2.1997
Epoch 5/10
239/239 382s 2s/step - accuracy: 0.2498 - loss: 2.5113 - val_accuracy: 0.3509 - val_loss: 2.0429
Epoch 6/10
239/239 383s 2s/step - accuracy: 0.2767 - loss: 2.3602 - val_accuracy: 0.3887 - val_loss: 1.9393
Epoch 7/10
239/239 382s 2s/step - accuracy: 0.3103 - loss: 2.2509 - val_accuracy: 0.4038 - val_loss: 1.8496
Epoch 8/10
239/239 382s 2s/step - accuracy: 0.3399 - loss: 2.1333 - val_accuracy: 0.4302 - val_loss: 1.7842
Epoch 9/10
239/239 381s 2s/step - accuracy: 0.3605 - loss: 2.0537 - val_accuracy: 0.4566 - val_loss: 1.7169
Epoch 10/10
239/239 382s 2s/step - accuracy: 0.3826 - loss: 1.9712 - val_accuracy: 0.4792 - val_loss: 1.6559
```

```
In [9]: plt.plot(history.history['loss'],label='Train_loss')
plt.plot(history.history['val_loss'],label='Val_loss')
plt.legend()
plt.xlabel('No of Epochs')
plt.ylabel('Loss')
plt.title('Loss vs Epochs')
plt.show()

plt.plot(history.history['accuracy'],label = 'Train_acc')
plt.plot(history.history['val_accuracy'],label = 'Val_acc')
plt.legend()
plt.xlabel('No. of Epochs')
plt.ylabel("Accuracy")
plt.title("Accuracy vs Epochs")
plt.show()
```





```
In [10]: test_data = ImageDataGenerator(rescale=1./255)
test_generator = test_data.flow_from_directory(test_path, target_size=(640, 640),
```

Found 265 images belonging to 53 classes.

```
In [11]: predictions = np.argmax(model.predict(test_generator), axis=1)
matrix = confusion_matrix(test_generator.labels, predictions)
report = classification_report(test_generator.labels, predictions, zero_division=1)
print("Classification Report:\n", report)
```

9/9 ————— 26s 2s/step

## Classification Report:

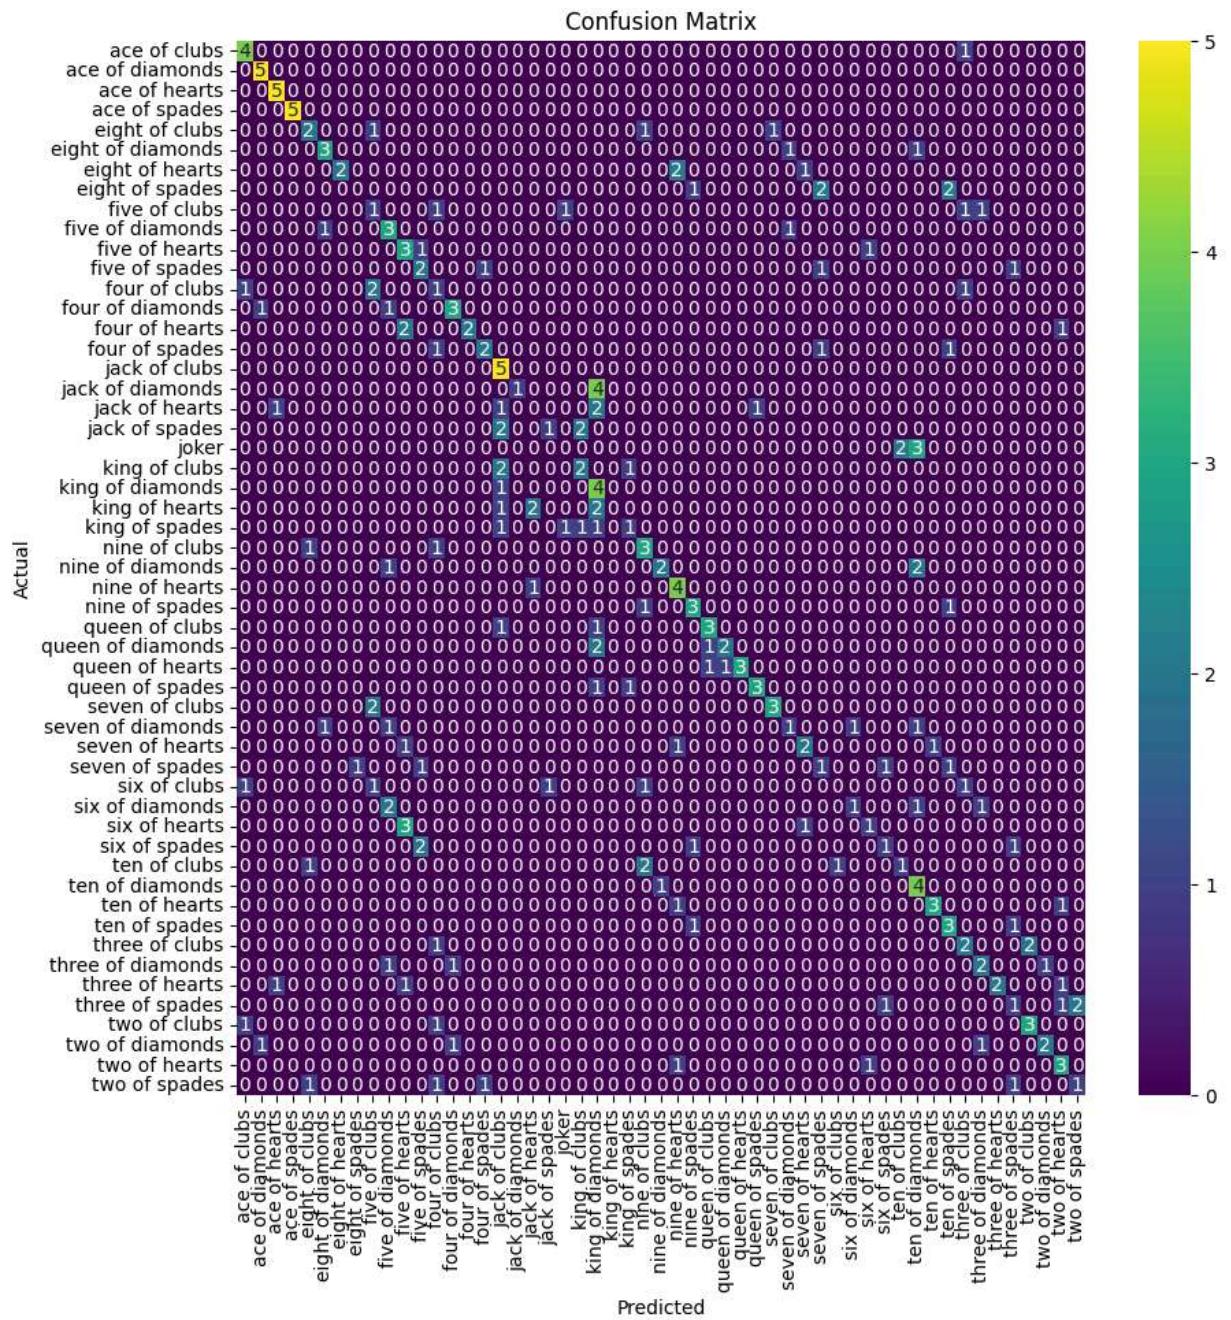
	precision	recall	f1-score	support
0	0.57	0.80	0.67	5
1	0.71	1.00	0.83	5
2	0.71	1.00	0.83	5
3	1.00	1.00	1.00	5
4	0.40	0.40	0.40	5
5	0.60	0.60	0.60	5
6	1.00	0.40	0.57	5
7	0.00	0.00	0.00	5
8	0.14	0.20	0.17	5
9	0.33	0.60	0.43	5
10	0.30	0.60	0.40	5
11	0.33	0.40	0.36	5
12	0.14	0.20	0.17	5
13	0.60	0.60	0.60	5
14	1.00	0.40	0.57	5
15	0.50	0.40	0.44	5
16	0.36	1.00	0.53	5
17	1.00	0.20	0.33	5
18	0.00	0.00	0.00	5
19	0.50	0.20	0.29	5
20	0.00	0.00	0.00	5
21	0.40	0.40	0.40	5
22	0.24	0.80	0.36	5
23	0.00	0.00	0.00	5
24	0.33	0.20	0.25	5
25	0.38	0.60	0.46	5
26	0.67	0.40	0.50	5
27	0.44	0.80	0.57	5
28	0.50	0.60	0.55	5
29	0.60	0.60	0.60	5
30	0.67	0.40	0.50	5
31	1.00	0.60	0.75	5
32	0.75	0.60	0.67	5
33	0.75	0.60	0.67	5
34	0.33	0.20	0.25	5
35	0.50	0.40	0.44	5
36	0.20	0.20	0.20	5
37	0.00	0.00	0.00	5
38	0.50	0.20	0.29	5
39	0.33	0.20	0.25	5
40	0.33	0.20	0.25	5
41	0.33	0.20	0.25	5
42	0.33	0.80	0.47	5
43	0.75	0.60	0.67	5
44	0.38	0.60	0.46	5
45	0.33	0.40	0.36	5
46	0.40	0.40	0.40	5
47	1.00	0.40	0.57	5
48	0.20	0.20	0.20	5
49	0.60	0.60	0.60	5
50	0.67	0.40	0.50	5
51	0.43	0.60	0.50	5
52	0.33	0.20	0.25	5

accuracy		0.44	265
macro avg	0.47	0.44	0.42
weighted avg	0.47	0.44	0.42

```
In [12]: pred = model.predict(test_generator,verbose=1)
test_generator_indices = np.argmax(pred,axis=1)
```

9/9 ━━━━━━━━ 13s 2s/step

```
In [14]: fig = plt.figure(figsize=(10,10))
sns.heatmap(matrix, annot=True, cmap='viridis')
plt.xticks(ticks=np.arange(53) + 0.5, labels=test_generator.class_indices, rotation=45)
plt.yticks(ticks=np.arange(53) + 0.5, labels=test_generator.class_indices, rotation=45)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
fig.savefig("Confusion Matrix",dpi=700)
```



In [ ]:

